

Lecture 9: Dataflow MoCs

Seyed-Hosein Attarzadeh-Niaki

Some Slides due to Edward Lee

Embedded Real-Time Systems

1

Review


- Composition of state machines
 - Synchronous composition
 - Asynchronous composition
- Hierarchical composition: StateCharts

Embedded Real-Time Systems

2

Dataflow Models

- Proven to be useful for specifying *streaming applications*
 - e.g., *signal processing* and *communications* domains
- *Simulation* of the algorithm at the functional or behavioral level
- *Synthesis* to software (e.g., a C program) or hardware (e.g., VHDL)
- Block-diagram based *visual programming*



3

Dataflow Signals (Streams)


- In dataflow signals, communication between the actors is done as *sequences of messages*
- Each message is called a *token*.
- A signal s is a *partial function* of the form

$$s: N \rightarrow V_s$$

V_s is the type of the signal
- The signal is defined on an initial segment $\{0, 1, \dots, n\} \subset N$, or (for infinite executions) on the entire set N .
- Unlike the synchronous reactive MoC, *no necessary time relation between inputs and outputs of an actor*.
 - They are *not synchronized*.

Dataflow Actors

- A (determinate) actor will be described as a function that **maps input sequences to output sequences**
- An **actor function** F , maps *entire* input sequences to *entire* output sequences
- A **firing function** f , maps a finite portion of the input sequences to output sequences
- A **firing rule** specifies the number of tokens required on each input port in order to fire the actor



$$F(x_1, x_2, x_3, \dots) = (ax_1, ax_2, ax_3, \dots)$$



$$\text{e.g., } f(x_1) = (ax_1)$$

Embedded Real-Time Systems

5

Composition of Dataflow Actors: Buffered Communication


- Since the firing of the actors is **asynchronous**, a token sent from one actor to another must be **buffered**.
- When the destination actor fires, it **consumes** one or more input tokens.
- Being able to execute a dataflow model forever is called an **unbounded execution**.
 - We need scheduling policies that deliver **bounded buffers**.
 - **Deadlock** is another challenge which appears in the cycles.




Embedded Real-Time Systems

6

Implementing Buffers for Dataflow Models



- Unbounded buffers require memory allocation and deallocation schemes.
- Bounded size buffers can be realized as *circular buffers* or *ring buffers*, in a statically allocated array.
 - A *read pointer* r is an index into the array referring to the first empty location. Increment this after each read.
 - A *fill count* n is unsigned number telling us how many data items are in the buffer. (or a *tail pointer*)
 - The next location to write to is $(r + n)$ modulo buffer length.
 - The buffer is empty if $n == 0$
 - The buffer is full if $n == \text{buffer length}$
 - Can implement n as a (counting) semaphore, providing mutual exclusion for code that changes n or r .



Embedded Real-Time Systems

7

Facts About (General) Dynamic Dataflow

- Whether there exists a schedule that does not *deadlock* is **undecidable**.
- Whether there exists a schedule that executes forever with *bounded memory* is **undecidable**.

Undecidable means that there is no algorithm that can answer the question in finite time for all finite models.

However, there are special cases where the models are analyzable

Embedded Real-Time Systems

8

Dataflow: many variants, still active area of research


- Computation graphs [Karp & Miller - 1966]
- Process networks [Kahn - 1974]
- Static dataflow [Dennis - 1974]
- Dynamic dataflow [Arvind, 1981]
- K-bounded loops [Culler, 1986]
- Synchronous dataflow [Lee & Messerschmitt, 1986]
- Structured dataflow [Kodosky, 1986] now
- PGM: Processing Graph Method [Kaplan, 1987]
- Synchronous languages [Lustre, Signal, 1980's]
- Well-behaved dataflow [Gao, 1992]
- Boolean dataflow [Buck and Lee, 1993]
- Multidimensional SDF [Lee, 1993]
- Cyclo-static dataflow [Lauwereins, 1994]
- Integer dataflow [Buck, 1994]
- Bounded dynamic dataflow [Lee and Parks, 1995]
- Heterochronous dataflow [Girault, Lee, & Lee, 1997]
- Parameterized dataflow [Bhattacharya and Bhattacharyya 2001]
- Structured dataflow (again) [Thies et al. 2002]
- ...

Lee 09: 6

Embedded Real-Time Systems

9

Synchronous Dataflow (SDF)




- The **number of tokens** consumed and produced by the firing of an actor is **constant**
- **Static analysis** can tell us whether we can schedule the firings to get a useful execution
- If so, then a **finite representation** of a schedule for such an execution can be created.

Embedded Real-Time Systems

10

Balance Equations



Let q_A, q_B be the number of firings of actors A and B.

Let p_C, c_C be the number of tokens produced and consumed on a connection C.


Then the system is *in balance* if for all connections C

$$q_A p_C = q_B c_C$$

where A produces tokens on C and B consumes them.

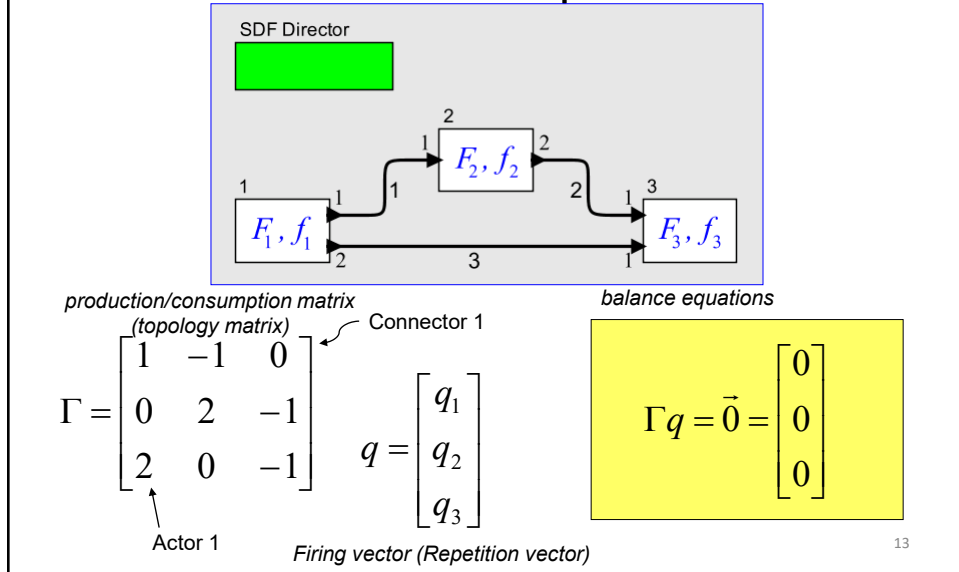
Example

Consider this example, where actors and arcs are numbered:




The balance equations imply that actor 3 must fire twice as often as the other two actors.

Compactly Representing The Balance Equations




Question

- What is the production/consumption matrix in this case?



$$\Gamma = \begin{bmatrix} 1 & 0 & -1 \\ 1 & -8 & 0 \end{bmatrix}$$

Example



A solution to the balance equations:


$$q = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \quad \Gamma = \begin{bmatrix} 1 & -1 & 0 \\ 0 & 2 & -1 \\ 2 & 0 & -1 \end{bmatrix} \quad \Gamma q = \vec{0}$$

This tells us that actor 3 must fire twice as often as actors 1 and 2.

Embedded Real-Time Systems

15

Example



But there are many solutions to the balance equations:

$$q = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} \quad q = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} \quad q = \begin{bmatrix} 2 \\ 2 \\ 4 \end{bmatrix} \quad q = \begin{bmatrix} -1 \\ -1 \\ -2 \end{bmatrix} \quad q = \begin{bmatrix} \pi \\ \pi \\ 2\pi \end{bmatrix} \quad \Gamma q = \vec{0}$$

For “well-behaved” models, there is a **unique least positive integer** solution.

Embedded Real-Time Systems

16

Least Positive Integer Solution to the Balance Equations

- Note that: if p_C, c_C , the number of tokens produced and consumed on a connection C, are **non-negative integers**, then the balance equation,

$$q_A p_C = q_B c_C$$

implies:

q_A is rational if and only if q_B is rational.

q_A is positive if and only if q_B is positive.

- Consequence:** Within any connected component, if there is any non-zero solution to the balance equations, then there is a unique least positive integer solution.

Rank of a Matrix

- The **rank** of a matrix Γ is the number of linearly independent rows or columns. The equation

$$\Gamma q = \bar{0}$$

is forming a linear combination of the columns of Γ .

- Such a linear combination can only yield the **zero vector** if the columns are **linearly dependent** (this is what it means to be *linearly dependent*).
- If Γ has a columns and b rows, the rank cannot exceed $\min(a, b)$.
- If the columns or rows of Γ are re-ordered, the resulting matrix has the same rank as Γ .

Rank of the Production/Consumption Matrix

- Let a be the number of actors in a connected graph. Then the *rank* of the production/consumption matrix Γ is $\leq a$. (why?)
- If the model is a *spanning tree* (meaning that there are barely enough connections to make it connected) then Γ has a rows and $a - 1$ columns.
- Theorem [Lee-Messerschmitt' 87]: Its rank is $a - 1$.
– Exercise: Prove it. (Hint: use induction).
- Corollary: the rank of any production/consumption matrix of a connected graph is either a or $a - 1$. (why?)

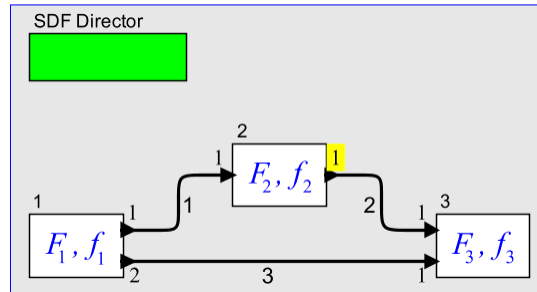
Consistent Models



- Let a be the number of actors in a connected model. The model is *consistent* if Γ has rank $a - 1$.
- If the rank is a , then the balance equations have only a *trivial solution* (zero firings).
- When Γ has rank $a - 1$, then the balance equations *always have a non-trivial solution*.

Example of an Inconsistent Model: No Non-Trivial Solution to the Balance Equations

$$\Gamma = \begin{bmatrix} 1 & -1 & 0 \\ 0 & 1 & -1 \\ 2 & 0 & -1 \end{bmatrix}$$



- This production/consumption matrix has rank 3, so there are no nontrivial solutions to the balance equations.
- Note that this model can execute forever, but it requires unbounded memory.

Embedded Real-Time Systems

21

Necessary and sufficient conditions


Consistency is a necessary condition to have a (bounded-memory) infinite execution.

Is it sufficient?

Embedded Real-Time Systems

22

Deadlock 1




Is this diagram consistent?

Embedded Real-Time Systems

23

Deadlock 2



Some dataflow models **cannot execute forever**. In the above model, the feedback loop injects **initial tokens**, but *not enough* for the model to execute.

Embedded Real-Time Systems

24

SDF: from static analysis to scheduling

- **Given:** SDF diagram
Find: a bounded-buffer schedule, if it exists
- **Step 0:** check whether diagram is **consistent**. If not, then no bounded-buffer schedule exists.
- **Step 1:** find an **integer solution** to $\Gamma q = 0$.
- **Step 2:** “decompose” the solution q into a schedule, making sure **buffers never become negative**.



Embedded Real-Time Systems

25

Step 2: “decomposing” the firing vector

Example 1:

Schedule = (1;2;3;3)




Embedded Real-Time Systems

26

Step 2: “decomposing” the firing vector

Example 2:



What happens if we try to run the previous procedure on this example?

So, we have both necessary and sufficient conditions for scheduling SDF graphs.


Embedded Real-Time Systems

27

A Key Question: If More Than One Actor is Fireable in Step 2, How do I Select One?

Optimization criteria that might be applied:

- Minimize buffer sizes.
- Minimize the number of actor activations.
- Minimize the size of the representation of the schedule (code size).




See S. S. Bhattacharyya, P. K. Murthy, and E. A. Lee, *Software Synthesis from Dataflow Graphs*, Kluwer Academic Press, 1996.

Beyond our scope here, but hints that it's an interesting problem...

Embedded Real-Time Systems

28

Minimum Buffer Schedule




ABABCABCABABCBCDEAFFFFFFBABCABCABABCDE
AFFFFFFBCABABCABABCBCDEAFFFFFFBCABABCABC
DEAFFFFFFBABCABCABABCBCDEAFFFFFFBABCABC
BABCBCDEAFFFFFFBCABABCABABCBCDEAFFFFFFBCA
FFFFFFBABCABCBCDEAFFFFFFBABCABCABABCBCDEAF
FFFFFFBABCABCABABCBCDEAFFFFFFBABCABCABABC
DEAFFFFFFBCABABCBCDEAFFFFFFBABCABCABABC
BCDEAFFFFFFBABCABCABABCBCDEAFFFFFFBCAFFFFFFB
ABCABCABABCBCDEAFFFFFFBCABABCBCDEAFFFFFFBA
BCABCABABCBCDEAFFFFFFBABCABCABABCBCDEAFFF
FFBCABABCABCABABCBCDEAFFFFFFBCABABCBCDEAF
FFFFFFBABCABCABABCBCDEAFFFFFFBBAFFFFFFBCABC
ABABCBCDEAFFFFFFBCABABCABABCBCDEAFFFFFFBCA
BABCBCBCDEAFFFFFFBABCABCABABCBCBCDEAFFFFFFB
ABCABCABABCBCDEAFFFFFFBCABABCABABCBCDEAF
FFFFFFBCABABCBCBCDEAFFFFFFBCA

Source: Shuvra Bhattacharyya

Embedded Real-Time Systems

29

Scheduling Tradeoffs (Bhattacharyya, Parks, Pino)



Scheduling strategy	Code	Data
Minimum buffer schedule, no looping	13735	32
Minimum buffer schedule, with looping	9400	32
Worst minimum code size schedule	170	1021
Best minimum code size schedule	170	264

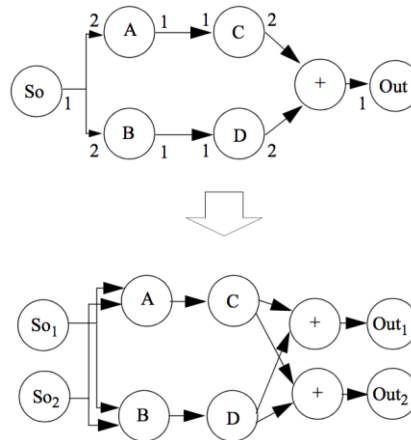
Source: Shuvra Bhattacharyya

Embedded Real-Time Systems

30

Homogeneous SDF

- An SDF graph in which every actor consumes and produces only one token from each of its inputs and outputs
- A general, consistent SDF graph that is not an HSDFG *can always be converted* into an equivalent HSDFG



Embedded Real-Time Systems

31

Graph-Based Techniques For (H)SDF Analysis

- Obtaining Feasible schedules using shortest paths in graphs
- Maximum Cycle Mean (MCM): Used to obtain the maximum available throughput of a self-timed HSDF graph ($\text{throughput}_{\max} = 1/\text{MCM}$).


$$MCM(G) = \max_{\text{cycle } C \text{ in } G} \left\{ \frac{\sum_{v \text{ is on } C} t(v)}{\text{Delay}(C)} \right\}$$

Embedded Real-Time Systems

32

Reflection

- How can we model limited buffer sizes in SDF?




Embedded Real-Time Systems

33

Parallel Scheduling of SDF Models

SDF is suitable for automated mapping onto parallel processors and synthesis of parallel circuits.



Many scheduling optimization problems can be formulated.



Sequential




Parallel

Embedded Real-Time Systems

34

Boolean Dataflow (BDF)

- SDF cannot directly express conditional firing
 - e.g., where an actor fires only if a token has a particular value.
- Two basic actors SWITCH and SELECT are added in boolean datflow
- Any Turing machine can be expressed as a BDF graph



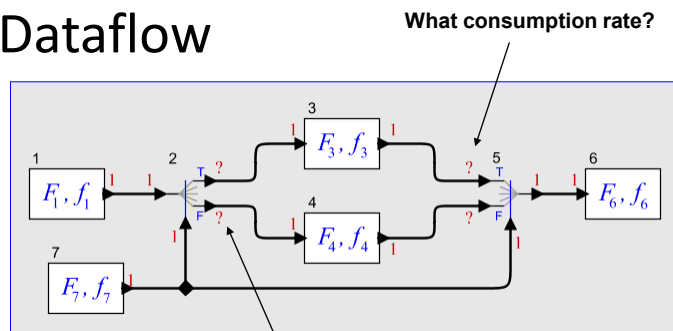
Embedded Real-Time Systems

35

Boolean Dataflow

Imperative equivalent:

```
while (true) {
  x = f1();
  b = f7();
  if (b) {
    y = f3(x);
  } else {
    y = f4(x);
  }
  f6(y);
}
```



What production rate?

The if-then-else model is not SDF.
But we can clearly give a bounded *quasi-static* schedule for it:
(1, 7, 2, b?3, !b?4, 5, 6)

guard

Embedded Real-Time Systems

36

