

Lecture 20: Real-Time Scheduling III

Seyed-Hosein Attarzadeh-Niaki

Based on the Slides by Edward Lee and Rodolfo Pellizzoni

Review

- Earliest Due Date (EDD) and Earliest Deadline First (EDF) scheduling
 - Optimality
- Precedence Constraints
 - Latest Deadline First (LDF) scheduling
 - EDF* scheduling

Outline

- Mutual exclusion
 - Priority inversion
 - Priority inheritance protocol
 - Deadlock
 - Priority ceiling protocol
- Aperiodic scheduling
 - Polling server
 - Sporadic server
- Multiprocessor scheduling
 - Brittleness
 - Richard's anomalies

Accounting for Mutual Exclusion

Recall from previous lectures:

- When threads access **shared resources**, they need to use mutexes to ensure **data integrity**.
- Mutexes can also **complicate** scheduling.

```

#include <pthread.h>
...
pthread_mutex_t lock;

void* addListener(notify listener) {
    pthread_mutex_lock(&lock);
    ...
    pthread_mutex_unlock(&lock);
}

void* update(int newValue) {
    pthread_mutex_lock(&lock);
    value = newValue;
    elementType* element = head;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
    pthread_mutex_unlock(&lock);
}

int main(void) {
    pthread_mutex_init(&lock, NULL);
    ...
}

```

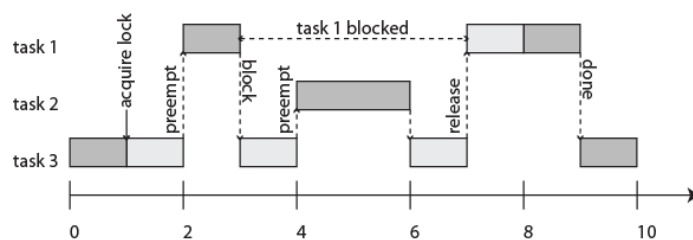
Recall mutual exclusion mechanism in pthreads

- Whenever a data structure is **shared** across threads, **access** to the data structure must usually be **atomic**.
- This is enforced using **mutexes**, or mutual exclusion locks.
- The code executed while holding a lock is called a **critical section**.

Embedded Real-Time Systems

5

Priority Inversion: A Hazard with Mutexes



- Task 1 has highest priority, task 3 lowest.
- Task 3 **acquires** a lock on a shared object, entering a **critical section**.
- It gets **preempted** by task 1, which then tries to acquire the lock and blocks.
- Task 2 **preempts** task 3 at time 4, keeping the higher priority task 1 blocked for an *unbounded amount of time*.
- In effect, the **priorities** of tasks 1 and 2 get **inverted**, since task 2 can keep task 1 waiting arbitrarily long.

Embedded Real-Time Systems

6

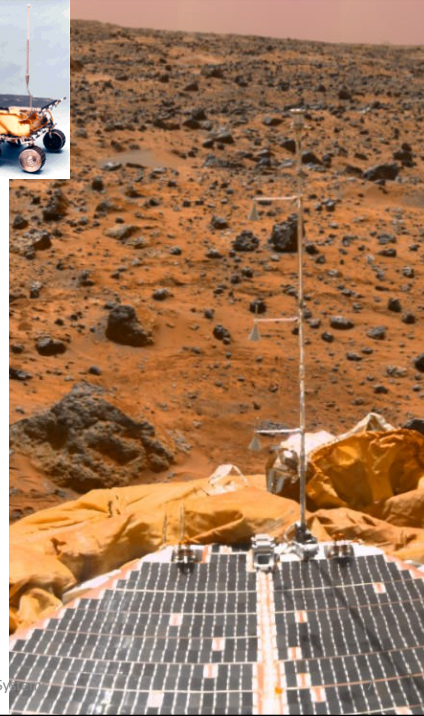
Mars Rover Pathfinder



- The Mars Rover Pathfinder landed on Mars on July 4th, 1997.
- A few days into the mission, the Pathfinder began **sporadically missing deadlines**, causing total system **resets**, each with **loss of data**.
- The problem was diagnosed on the ground as **priority inversion**, where a low priority meteorological task was holding a lock **blocking a high-priority task** while medium priority tasks executed.

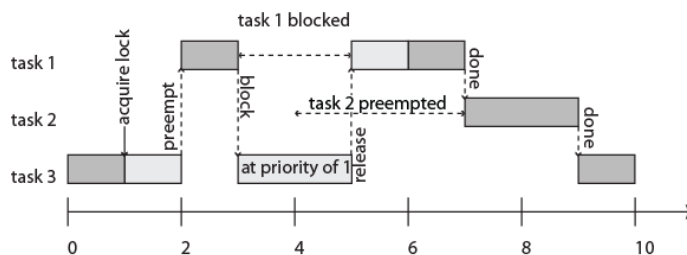
Source: RISKS-19.49 on the comp.programming.threads newsgroup, December 07, 1997, by Mike Jones (mbj@MICROSOFT.com).

Embedded Real-Time Systems



Priority Inheritance Protocol (PIP)

(Sha, Rajkumar, Lehoczky, 1990)



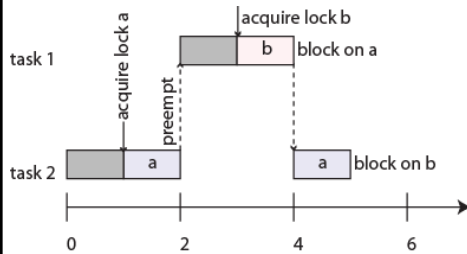
➤ The task that holds the lock **inherits** the priority of the blocked task.

- Task 1 has highest priority, task 3 lowest.
- Task 3 acquires a lock on a shared object, entering a **critical section**.
- It gets **preempted** by task 1, which then tries to acquire the lock and blocks.
- Task 3 inherits the priority of task 1, preventing preemption by task 2.

Embedded Real-Time Systems

8

Deadlock



- The **lower priority** task starts first and **acquires lock a**, then gets **preempted** by the **higher priority** task, which **acquires lock b** and then **blocks** trying to **acquire lock a**.
- The **lower priority** task then **blocks** trying to **acquire lock b**, and **no further progress is possible**.

```
#include <pthread.h>
...
pthread_mutex_t lock_a, lock_b;

void* thread_1_function(void* arg) {
    pthread_mutex_lock(&lock_b);
    ...
    pthread_mutex_lock(&lock_a);
    ...
    pthread_mutex_unlock(&lock_a);
    ...
    pthread_mutex_unlock(&lock_b);
    ...
}

void* thread_2_function(void* arg) {
    pthread_mutex_lock(&lock_a);
    ...
    pthread_mutex_lock(&lock_b);
    ...
    pthread_mutex_unlock(&lock_b);
    ...
    pthread_mutex_unlock(&lock_a);
    ...
}
```

Embedded Real-Time Systems

9

Priority Ceiling Protocol (PCP)

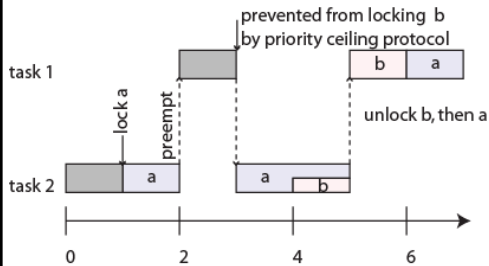
(Sha, Rajkumar, Lehoczky, 1990)

- Every lock or semaphore is assigned a **priority ceiling** equal to the priority of the *highest-priority task that can lock it*.
 - Can one automatically compute the priority ceiling?
- A task T can acquire a lock only if the task's priority is *strictly higher than the priority ceilings of all locks currently held by other tasks*.
 - Intuition: the task T will not later try to acquire these locks held by other tasks
 - Locks that are not held by any task don't affect the task
- This **prevents deadlocks**.
- There are **extensions** supporting **dynamic priorities** and **dynamic creations of locks** (stack resource policy).

Embedded Real-Time Systems

10

Priority Ceiling Protocol



In this version, locks a and b have **priority ceilings** equal to the priority of task 1.

At time 3, task 1 attempts to lock b, but **it can't** because task 2 currently holds lock a, which has priority ceiling equal to the priority of task 1.

```
#include <pthread.h>
...
pthread_mutex_t lock_a, lock_b;

void* thread_1_function(void* arg) {
    pthread_mutex_lock(&lock_b);
    ...
    pthread_mutex_lock(&lock_a);
    ...
    pthread_mutex_unlock(&lock_a);
    ...
    pthread_mutex_unlock(&lock_b);
    ...
}

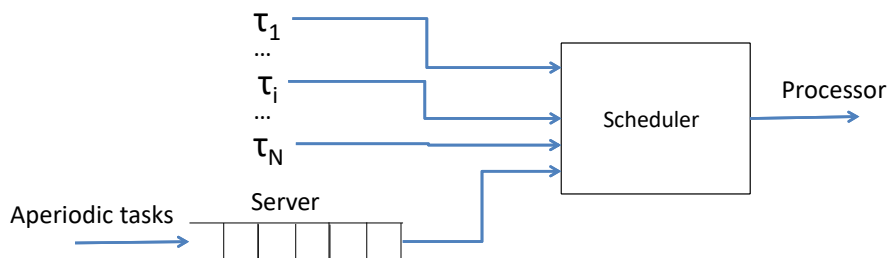
void* thread_2_function(void* arg) {
    pthread_mutex_lock(&lock_a);
    ...
    pthread_mutex_lock(&lock_b);
    ...
    pthread_mutex_unlock(&lock_b);
    ...
    pthread_mutex_unlock(&lock_a);
    ...
}
```

Embedded Real-Time Systems

11

Aperiodic Tasks

- What happens if we **mix aperiodic and periodic** tasks?
- Main idea:** ensure that *periodic tasks remain schedulable* no matter what.
- Aperiodic server**
 - Insert aperiodic tasks into a queue (server)
 - Scheduler picks among periodic tasks or the aperiodic server
- Server Goals**
 - Minimize response time of aperiodic tasks
 - Low overhead



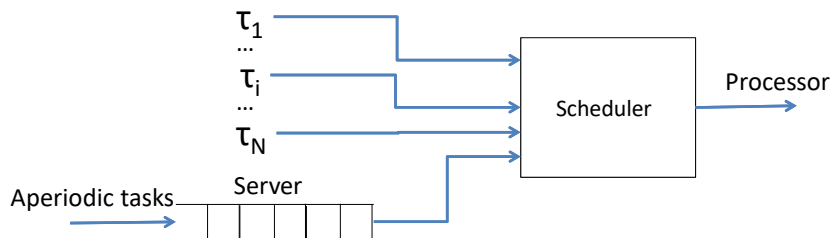
Embedded Real-Time Systems

12

Aperiodic Servers

Solution#1: background server

- Execute aperiodics whenever the CPU is not running a periodic task (i.e., the **server has lowest priority**)
- Problem: response time can be very **high**.



Embedded Real-Time Systems

13

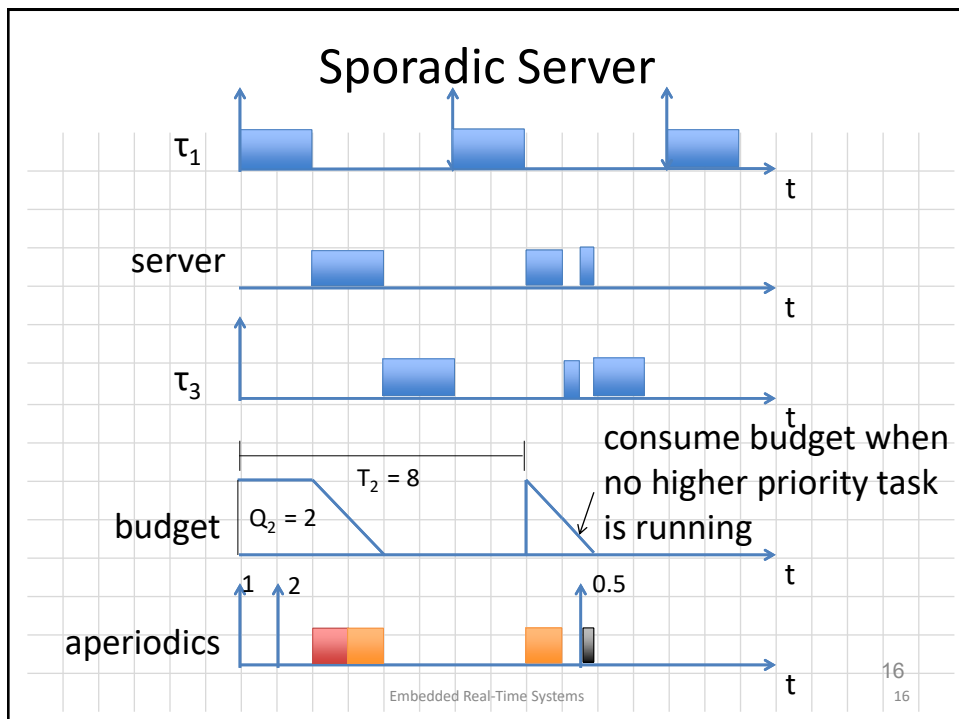
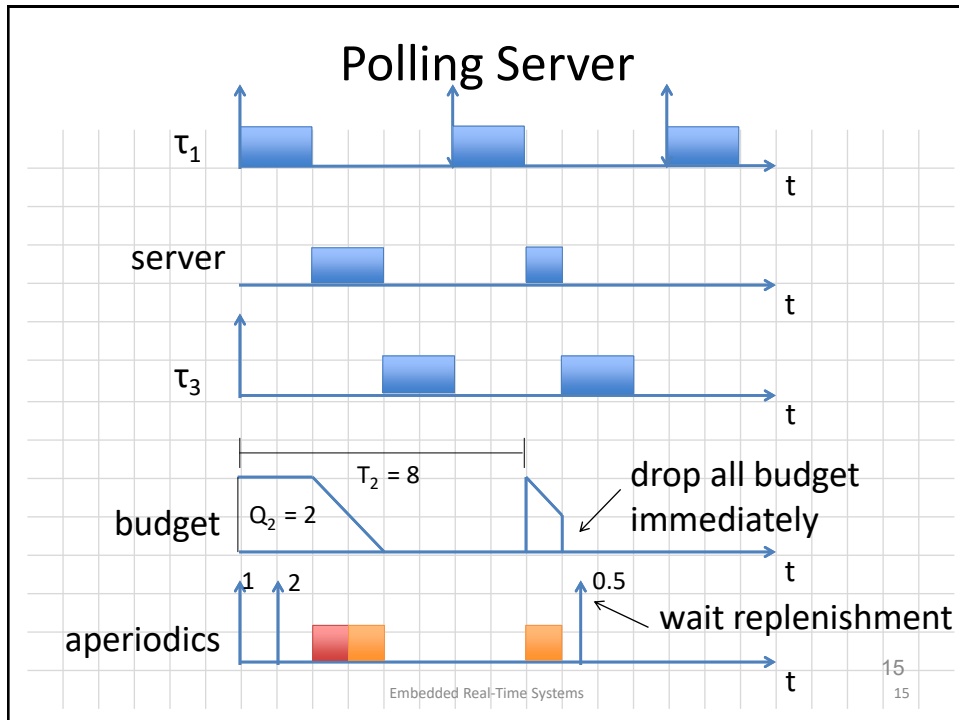
Aperiodic Servers

Solution#2: budget-based server

- Server is assigned **budget Q_i** , **period T_i** .
- The server behaves like a periodic task with $C_i = Q_i$ and period T_i .
- When the scheduler picks the server, if there is budget left, the server executes an aperiodic in the queue consuming its budget.
- When budget=0, server waits until next period, then **replenish** budget to Q_i .
- **Problem**: what happens if the scheduler picks the server and there are no queued aperiodic tasks?
 - “**Dumb**” **servers** (polling server) lose budget.
 - “**Smart**” **servers** (ex: sporadic server) keep the budget but modify their activation (recharge) time.

Embedded Real-Time Systems

14



Scheduling on Multiprocessor

Solution #1: partitioning

- **Statically assign** tasks among M processors.
- Ex: EDF. Each core is schedulable if sum of utilizations of tasks assigned to that core ≤ 1 .
- Problem can be rephrased as: given a set of objects with known sizes (task utilizations), place them into M equal-size containers.
 - Classic bin-packing problem
 - NP-hard

Solution #2: global scheduling

- Keep a global scheduling queue.
- Whenever there is a free core, pick one task from the queue and schedule it on the core.

Embedded Real-Time Systems

17

Scheduling on Multiprocessor

- Typical scheduling goal: **minimizing makespan**.
- In practice, real-time adoption of multiprocessor is limited, especially for hard (real-time) systems.
- **Partitioned scheduling** preferred.
- Three *issues with global scheduling*
 1. Increases **unpredictability** – tasks can migrate among cores.
 2. Much more **complex** to implement.
 3. Does **not necessarily perform better** than partitioned.

Embedded Real-Time Systems

18

Global Scheduling Strategies

- Hu level scheduling algorithm

- Assigns a priority to each task τ based on the *level*

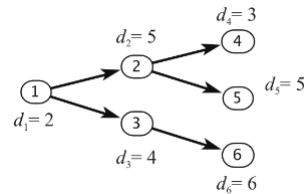
- Largest sum of exec times of tasks on paths from τ to leaf tasks.

- Larger level \rightarrow higher priority

- **Critical path-based**

- Example priorities:

- High τ_1 ; Medium τ_2, τ_3 ; Low τ_4, τ_5, τ_6



- List scheduler

- Sorts the tasks by priorities,

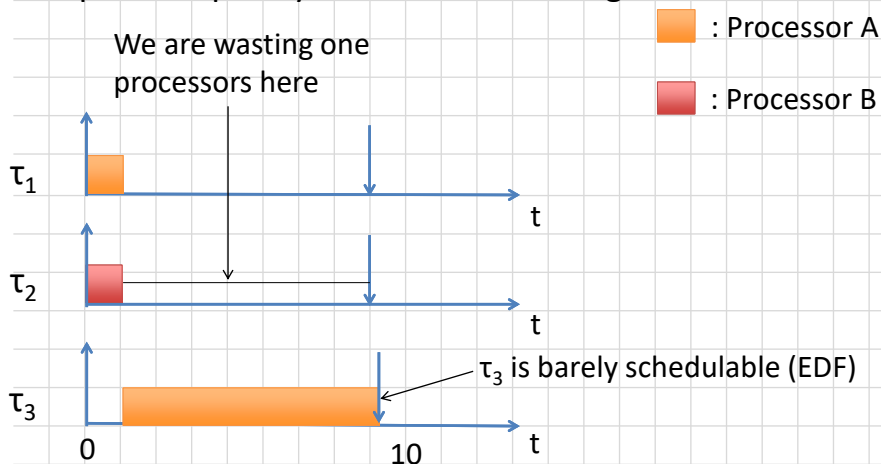
- assigns them to processors in the order of the sorted list as processors become available.

Embedded Real-Time Systems

19

Improving Global Scheduling

- Problem: both fixed-priority and EDF scheduling perform poorly when there are long tasks.

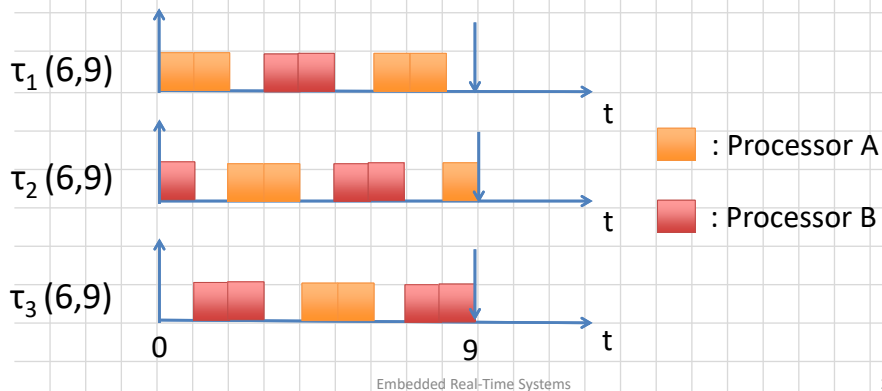


Embedded Real-Time Systems

20

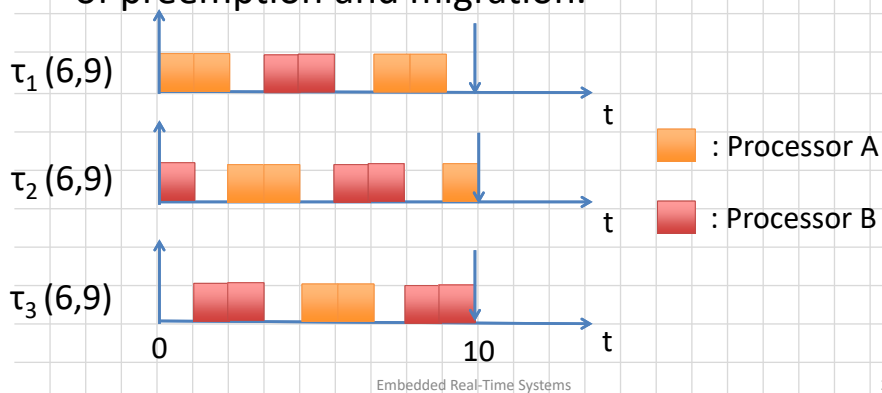
Improving Global Scheduling

- Optimal algorithm: p-fair.
- Split the tasks into small chunks.
- Allocate them on the cores in a “fair” way.



Improving Global Scheduling

- Pros: task set schedulable on M cores iff $U \leq M$; the algorithm is optimal.
- Cons: this does not take into account the cost of preemption and migration.



Scheduling on Multiprocessor

Some other details...

- There are several (only sufficient) schedulability analyses for EDF and FP – both based on utilization bounds and response time...
- There are extensions for resource sharing protocols and aperiodic servers to multicores...
- Very active research topic, but often ignores the main problem – how do we determine the worst-case computation time?

Embedded Real-Time Systems

23

Brittleness

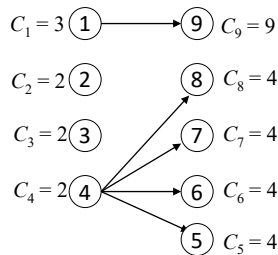
- In general, all thread scheduling algorithms are **brittle**: small changes can have big, unexpected consequences.
- A good illustration of this is with multiprocessor (or multicore) schedules.

Theorem (Richard Graham, 1976): If a task set with fixed priorities, execution times, and precedence constraints is scheduled according to priorities on a fixed number of processors, then **increasing the number of processors**, **reducing execution times**, or **weakening precedence constraints** can **increase** the schedule length.

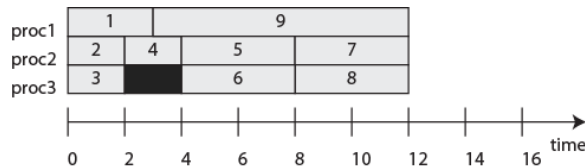
Embedded Real-Time Systems

24

Richard's Anomalies



9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:

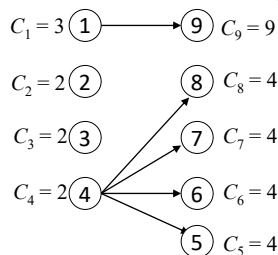


What happens if you increase the number of processors to four?

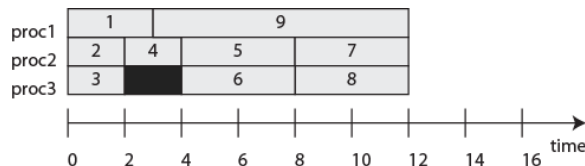
Embedded Real-Time Systems

25

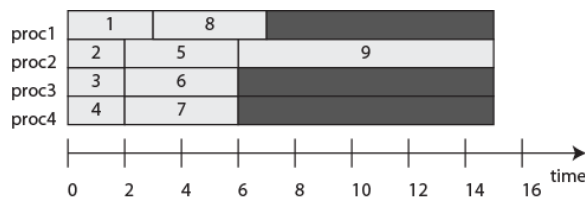
Richard's Anomalies: Increasing the number of processors



9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:



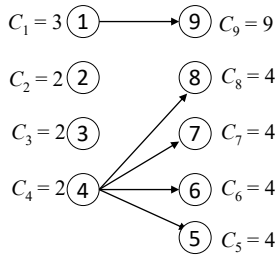
The priority-based schedule with four processors has a longer execution time.



Embedded Real-Time Systems

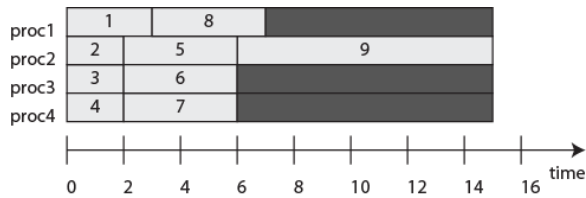
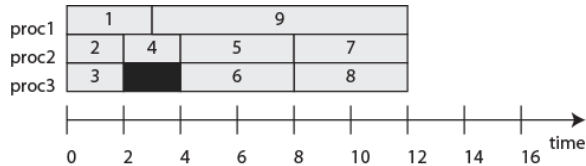
26

Greedy Scheduling



Priority-based scheduling is “greedy.” A smarter scheduler for this example could hold off scheduling 5, 6, or 7, leaving a processor idle for one time unit.

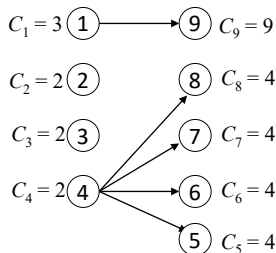
9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:



Embedded Real-Time Systems

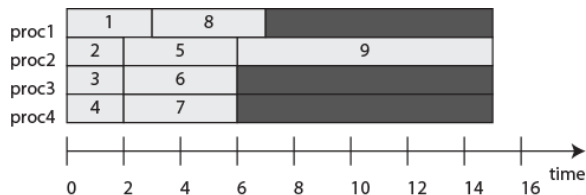
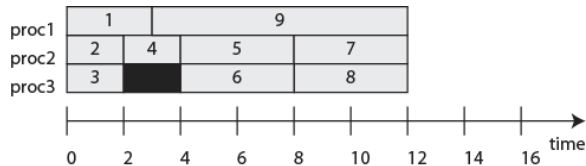
27

Greedy scheduling may be the only practical option.



If tasks “arrive” (become known to the scheduler) only after their predecessor completes, then greedy scheduling may be the only practical option.

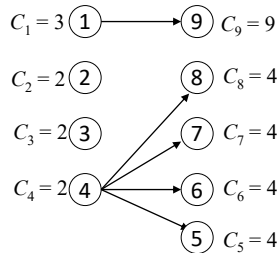
9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:



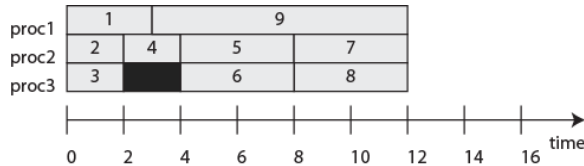
Embedded Real-Time Systems

28

Richard's Anomalies



9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:

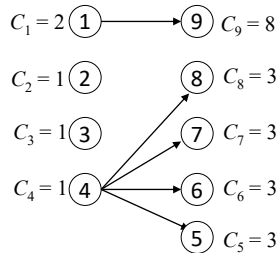


What happens if you reduce all computation times by 1?

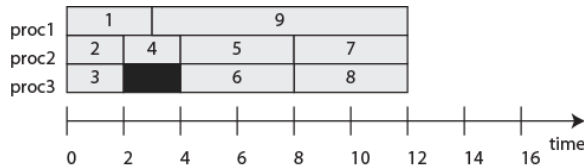
Embedded Real-Time Systems

29

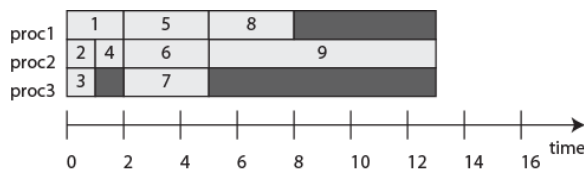
Richard's Anomalies: Reducing computation times



9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:



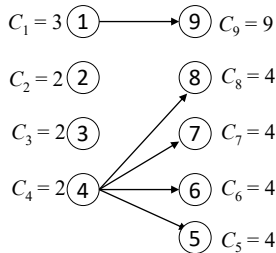
Reducing the computation times by 1 also results in a longer execution time.



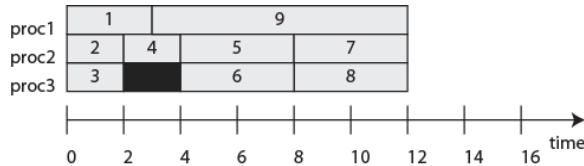
Embedded Real-Time Systems

30

Richard's Anomalies



9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:

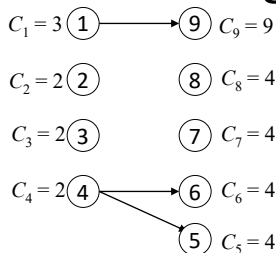


What happens if you remove the precedence constraints (4,8) and (4,7)?

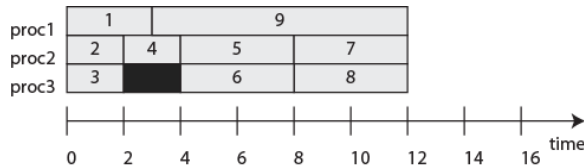
Embedded Real-Time Systems

31

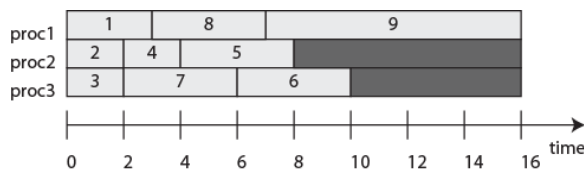
Richard's Anomalies: Weakening the precedence constraints



9 tasks with precedences and the shown execution times, where lower numbered tasks have higher priority than higher numbered tasks. Priority-based 3 processor schedule:



Weakening precedence constraints can also result in a longer schedule.

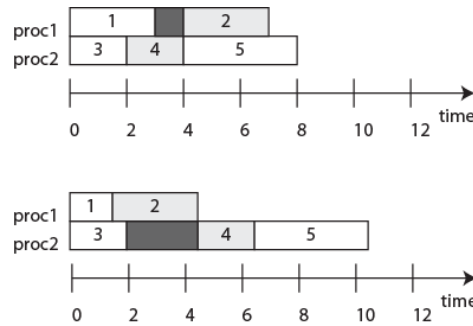


Embedded Real-Time Systems

32

Richard's Anomalies with Mutexes: Reducing Execution Time

- Assume tasks 2 and 4 share the same resource in exclusive mode, and tasks are statically allocated to processors.
- If the execution time of task 1 is reduced, the schedule length increases:



Embedded Real-Time Systems

33

In short...

- Timing behavior** under all known task scheduling strategies is **brittle**.
 - Small changes can have big (and unexpected) consequences.
- Unfortunately, since execution times are so hard to predict, such brittleness can result in **unexpected system failures**.

Embedded Real-Time Systems

34